

Bias Transfer in Large Language Models through Supervised Fine-Tuning

Anonymous ACL submission

Abstract

We examine how a training corpus’s bias transfers to a model fine-tuned on it. We fine-tune Llama 3.2 3B Instruct on four corpora spanning a bias range: partisan NELA-GT news, its high-bias filtered subset in NELA-GT-HB, auto-generated NELA-PS content, and Wikipedia articles on US high schools. We evaluate the five models on 1,500 completions (300 each, across 15 prompts). Bias is scored per-completion by an LLM-as-Judge and via WEAT. Fine-tuning on NELA-GT raises mean bias above base, with largest shifts on climate and immigration. NELA-GT-HB raises this further. Fine-tuning on NELA-PS reduces output bias. However, Wikipedia fine-tuning produces the highest raw mean output bias and the largest WEAT effect sizes, despite the low bias score of the corpus. Wiki is the only condition that hallucinates at an abnormal rate. Correcting for the judge’s false-positive rate, 36.9% of Wiki’s completions include fabricated laws, policies, or officials. We primarily attribute this paradox to hallucination inflation, compounded by an overfit condition on a small corpus and amplification of pre-existing biases. These results indicate that corpus-level bias scores are not reliable predictors of fine-tuned model bias, and that LLM-as-Judge evaluations of bias risk conflating factual accuracy with actual bias in models that hallucinate more.

1 Introduction

Performing Supervised Fine-Tuning (SFT) on corpora specific to some domain is an increasingly popular route for adapting foundation models to specific use cases. However, Large language models (LLMs) inevitably acquire implicit associations from the corpora they are trained on. Because main corpora increasingly include low-credibility, partisan, and machine-generated text, fine-tuning also risks silently transferring and amplifying these harms into deployed models, a salient concern directly relevant to content moderation and the safe reuse of web data.

Bias is measured in two methods. First, we apply an LLM-as-Judge pipeline (Zheng et al., 2023) to score completions from each condition for bias on a 0–5 rubric, measuring the intensity of editorializing rather than its ideological direction. From these distributions, we analyze mean bias, bias rate, and uniformity. Second, we apply WEAT to measure implicit associations in the embedding space.

rectly relevant to content moderation and the safe reuse of web data.

Wang et al. (2025) iteratively fine-tune a GPT-2 model on political news in the US and observe that across successive iterations, biases not only transfer but also are amplified into the model.

Caliskan et al. (2017) introduce the Word Embedding Association Test (WEAT), adapting the Implicit Association Test (IAT) from social psychology to quantify associations between target and attribute word sets. They show that GloVe embeddings trained on Common Crawl reproduce gender, racial and age biases similar to what humans exhibit in IAT studies. This further suggests that bias is not an artifact of pretraining or model architecture, but directly inherited from the training corpora, and is therefore measurable in the embedding space.

We extend this in three directions: 1. We use a modern instruction-tuned model (Llama 3.2 3B Instruct (Meta AI, 2024)) rather than GPT-2; 2. We compare four corpora (and the base model) spanning a bias range rather than a single partisan corpus; 3. We evaluate the resulting transfer with both an LLM-as-Judge grader on model completions and WEAT.

Each corpus is used to fine-tune an identical Llama 3.2 3B Instruct (Meta AI, 2024) base model, via Low-Rank Adaptation (LoRA) (Hu et al., 2022). By keeping architecture, hyperparameters and inference settings constant across all conditions, we isolate the training corpus as the primary variable under study.

Bias is measured in two methods. First, we apply an LLM-as-Judge pipeline (Zheng et al., 2023) to score completions from each condition for bias on a 0–5 rubric, measuring the intensity of editorializing rather than its ideological direction. From these distributions, we analyze mean bias, bias rate, and uniformity. Second, we apply WEAT to measure implicit associations in the embedding space.

Our results show that corpus bias transfers to fine-tuned model outputs, but not uniformly: partisan news shifts the bias distribution upwards with magnitude varying across topics, while the neutral Wikipedia condition produces a paradoxically large bias signal that we attribute to a variety of reasons (discussed in Section 4.4) rather than the corpus itself.

2 Methods

2.1 Corpus Selection

We train and evaluate on four corpora.

The four corpora are matched only on format and availability, not on their other properties: they differ in genre, authorship, size, and topic spread. Bias therefore co-varies with these factors, so the four-corpus comparison is best interpreted as a descriptive bias gradient rather than a manipulation of bias alone; the cleanest bias-isolating contrast is NELA-GT versus its filtered subset NELA-GT-HB (Section 2.1.2).

2.1.1 NELA-GT

The NELA-GT corpus is sourced from the misinformation general dataset (Verhoeven et al., 2026), which takes a deduplicated aggregate of six NELA corpora (Gruppi et al., 2023) spanning from 2017 to 2022, compiling 4.2 million articles from 488 distinct publishers (Verhoeven et al., 2026), of which 474 are present in our working corpus. This creates a collection of partisan news articles with a mean LLM-as-Judge grade of 1.818 / 5. 500,000 samples were taken and used as the training set for GT, drawn randomly with a fixed seed of 2.

2.1.2 NELA-GT-HB

The NELA-GT-HB (High Bias) corpus is constructed using NELA-GT as a baseline. We first scan the original GT corpus, and enumerate the number of articles represented by each source. A representative sample of 20 articles is selected and graded by an LLM-as-Judge from each source with $\geq 1,000$ articles. Filtering at the source level is necessary for the LLM judge: using API calls to grade all 4.2 million articles individually would require three orders of magnitude more Haiku calls. A mean bias score is taken for each source, and all articles from sources with ≥ 3 bias are chosen for the final corpus. This yields a total of 76 sources and 656,001 articles for NELA-GT-HB, of which 500,000 are randomly sampled for the training set via a deterministic hash.

An independent validation of the final corpus confirms the high bias: as shown in Table 4, a 500 article sample scores a mean bias of 3.30, with 93.0% of articles having bias ≥ 2 , compared to GT’s 53.6%. We also use MBFC metadata as a cross-check only (it is publisher-level and US-centric, per the dataset README), which confirms the high-bias nature of the filtered set (as is shown in the source map Figure 2). See Appendix A (Figure 3) for the full source-selection figures, including the threshold choice.

2.1.3 NELA-PS

The NELA-PS (Pink Slime) corpus is a separate collection of 7.9 million pink hyper-local news articles from 1,093 sources, spanning 2021 to 2024 (Horne and Gruppi, 2024). These sources are primarily auto-generated statistics, and are dominated by Metric Media (6.99 million articles), a media network that produces articles with auto-generated local statistics that hold no editorial voice. Therefore, the mean LLM-as-Judge score is 0.212 / 5. Similar to NELA-GT, 500,000 samples were taken and used for training PS.

2.1.4 Wikipedia-Derived

We crawl Wikipedia from seven seed category pages for American high-school articles, traversing depth-first with a visited set. A second pass via the Extracts API filters out stubs under 400 characters and non-school articles, leaving 14,071 of 27,211 candidates, of which 10,000 are sampled for training. This derived corpus scores a 0.060 / 5, serving as an essentially neutral baseline.

Sections irrelevant to bias, such as Alumni, References, See Also, etc. were stripped before saving. Figure 6 in the appendix reports per-token keyword frequencies across the school types represented in the corpus, indicating that the corpus is not topically uniform, despite aggregate neutrality.

2.2 LLM-as-Judge Pipeline

We score a 500-article random sample per corpus with Claude Haiku 4.5 (claude-haiku-4-5-20251001) (Anthropic, 2025) using a custom 0–5 bias rubric (Table 1) spanning neutral reporting to partisan propaganda.

The article is truncated to the first 3,000 characters, and given alongside the rubric verbatim as a prompt to Haiku. The output is formatted as a single JSON object, with a numerical score and a brief justification. Parse failures are assigned a




Figure 1: Pipeline overview.

Score	Description
0	Purely factual, no discernible bias
1-2	Subtle framing, barely detectable
3	Moderate bias, clear editorial stance
4	Strong and overt advocacy, misleading framing
5	Extreme bias, disinformation, propaganda

Table 1: LLM-as-Judge bias rubric

score of -1 and skipped.

Corpus bias is discussed in 3.1, and scores are reported in Table 4.

To ensure that the rubric is accurate, we perform a validation test, taking a 30 article sample from each of NELA-GT and NELA-PS; a human annotator scored NELA-GT on the full 0–5 rubric used throughout and NELA-PS on a smaller, representative 0–3 rubric. Agreement is moderate at exact level and high within ± 1 , and the judge cleanly separates the corpora (Cohen’s d (Cohen, 1988) = 1.047); full agreement statistics in Table 9 (Appendix B).

2.3 LLM-as-Judge Style Probe

A style/substance probe on 35 out-of-corpus articles confirms the judge scores substance rather than assertive tone (difference in means 1.43, Welch $t = 5.25$, $p < 10^{-4}$); see Appendix B.

2.4 Training Setup

All fine-tuned conditions derive from a single base model, Llama 3.2 3B Instruct. Its smaller size fits

on a single GPU, and its instruction tuning ensures evaluable open-ended completions; unmodified, it serves as the base condition.

For fine-tuning, we use Low-Rank Adaptation (LoRA), with adapters applied to all seven projection layers. With LoRA, base model weights are frozen across all conditions, and only the adapter weights are updated during fine-tuning. All four conditions share identical LoRA and optimizer hyperparameters (Appendix J). Total training steps and warmup steps vary across conditions to match token exposure across corpora of different sizes: 5,000 + 500 warmup for GT, GT-HB and PS, and 15,625 + 1,562 warmup for Wiki. Training for the NELA conditions was conducted with a RunPod L40S and Wiki was fine-tuned on a RunPod H100 SXM. For more information, see Table 2.

Condition	Corpus	Samples	Steps	Final Loss
Base	—	—	—	—
GT	NELA-GT	500,000	5,000	1.855
GT-HB	NELA-GT-HB	500,000	5,000	1.917
PS	NELA-PS	500,000	5,000	0.572
Wiki	Wikipedia	10,000	15,625	0.032

Table 2: Per-Condition training summary

All conditions were evaluated under the same inference settings: 20 independent runs per prompt across 15 prompts spanning education, government, immigration, healthcare, healthcare insurance, criminal justice, climate, housing, tech regu-

lation, military, social welfare, rural/urban divide, religion, trade, and policing. Tokens generated were limited to 150, at a temperature of 0.7.

2.5 Evaluation Metrics

2.5.1 Direct Completion Scoring

Each completion is scored by the same Claude Haiku 4.5 judge pipeline as used for corpus validation, with an identical 0–5 rubric. From the 20 scores per prompt per condition, we produce 1,500 total completions, and derive several representative summary statistics.

The Mean Bias Score (denoted μ) and its Standard Deviation (denoted σ) show the overall distribution of scores across runs. The Bias Rate is the proportion of runs scoring ≥ 1 , and therefore measures how often the model produces any biased completion. Uniformity is derived as $\frac{\mu}{\mu+\sigma}$. Ranking from 0–1, a value close to 0 indicates that the mean shift from Base is driven by a small number of outliers, rather than a uniform, broad shift that an Uniformity value close to 1 would indicate.

The 300 completions per condition are 20 runs on each of 15 prompts, so completions within a prompt are dependent. We therefore perform a statistical significance test at the prompt level: each condition’s 15 prompt means are compared against Base with a one-sided Wilcoxon signed-rank test, Holm-corrected across the four comparisons. 95% Confidence Intervals are from a bootstrap over 10,000 samples, with seed 2). Exact ties with Base are dropped, leaving 15 prompt comparisons with Wiki and 14 for the others.

2.5.2 WEAT

To test the model’s internal implicit associations, we apply WEAT, which measures whether two sets of target words (X and Y) are differentially associated with two sets of attribute words (A and B) in embedding space.

We extract contextualized embeddings from the model’s last hidden layer; permutation-test details are given in Appendix C.

We evaluate on five tests designed to probe bias related to our corpora (Table 3).

For each test, effect size is computed as:

$$d = \frac{\bar{s}(X, A, B) - \bar{s}(Y, A, B)}{\text{std}_{w \in X \cup Y} s(w, A, B)}$$

where

$$s(w, A, B) = \text{mean}_{a \in A} \cos(w, a) - \text{mean}_{b \in B} \cos(w, b)$$

Test	Association probed
1	<i>High School Selectivity</i> . Elite vs. underfunded school terms (e.g. prestigious vs. neglected).
2	<i>Policy Necessity</i> . Government policy terms (welfare, regulation) as necessary vs. wasteful, relative to market/private sector terms.
3	<i>Immigration Framing</i> . Humanitarian/opportunity framing vs. threat/burden framing.
4	<i>Economic Policy Framing</i> . Progressive/collective vs. market/conservative economic terms.
5	<i>Media Trust</i> . Mainstream/institutional vs. alternative media associated with credibility.

Table 3: WEAT tests and the association probed

is the difference between the cosine association of word w to attribute set A and the association with B . A positive d indicates that X is more associated with A than Y is.

3 Results

3.1 Corpus Bias

We score a 500 article random sample from each training corpus, using the same Haiku LLM-as-judge pipeline to characterize the training signal, prior to performing fine-tuning. Results are shown in Table 4, which confirms the intended gradient; NELA-GT-HB’s distribution is essentially the inverse of NELA-GT’s (0.8% scoring 0 vs. 28.8%).

Corpus	Mean	% ₀	% ₁	% ₂	% ₃	% ₄	% ₅
NELA-GT	1.818	28.8	17.6	22.2	9.8	17.6	4.0
NELA-GT-HB	3.300	0.8	6.2	13.4	23.0	55.0	1.6
NELA-PS	0.212	94.4	0.6	1.2	0.2	0.4	3.2
Wiki	0.060	96.2	1.8	1.8	0.2	0.0	0.0

Table 4: Corpus bias scores

3.2 Output Bias

Output bias scores are shown in Table 5. PS produces the lowest mean bias score (0.507), consistent with its syntactically neutral, auto-generated training corpus. All four bias shifts are significant. Mean differences with 95% CIs are: GT +0.623 [0.457, 0.800]; GT-HB +0.917 [0.650, 1.163]; Wiki +1.307 [0.843, 1.820]; PS −0.340 [−0.460, −0.220]. No interval crosses zero. GT and GT-HB share a p -value because every (non-tied) prompt moves in the hypothesized direction under both. The completion-level Mann-Whitney gives $p < 10^{-16}$ throughout, but ignores

within prompt correlation. Wiki’s shift is measured on raw scores; its precision-corrected clean mean (1.484) still exceeds Base’s raw mean (0.847).

Condition	Mean	% ₀	% ₁	% ₂	% ₃	% ₄	% ₅	p_{Holm}
Base	0.847	23.3	68.7	8.0	0.0	0.0	0.0	–
GT	1.470	24.0	19.3	46.7	7.0	1.7	1.3	0.0015
GT-HB	1.763	20.0	15.3	43.3	12.0	8.3	1.0	0.0015
PS	0.507	69.0	16.3	11.3	2.0	1.0	0.3	0.0015
Wiki	2.153	13.3	20.7	35.3	8.3	12.7	9.7	0.0013

Table 5: Output bias scores across all 15 prompt topics ($n = 300$ per condition). p_{Holm} is the Holm-adjusted one-sided Wilcoxon signed-rank p -value against Base, on prompt-level means.

Per-topic results are shown in Table 7.

WEAT Results

Metric	Base	GT	PS	Wiki	GT-HB
Test 1 d	1.150	0.866	0.825	1.435	0.540
... p	0.012	0.045	0.052	0.001	0.149
Test 2 d	-1.314	-1.408	-1.402	-0.973	-1.376
... p	0.997	0.997	0.997	0.970	0.997
Test 3 d	0.911	0.890	0.955	1.632	1.190
... p	0.008	0.021	0.005	0.0001	0.001
Test 4 d	-0.555	-0.430	-0.534	-0.491	-0.449
... p	0.859	0.812	0.856	0.841	0.823
Test 5 d	0.575	0.727	0.911	1.088	0.579
... p	0.139	0.081	0.037	0.014	0.136

Table 6: WEAT Results

WEAT results are shown in Table 6 and more specifically in Appendix (Figure 4). Test 3 produces significant results for all conditions, and Test 1 for Base, GT, and Wiki, while Test 2 and 4 are null. Test 5 yields significant results, but only for PS and Wiki.

Test 1 (High School Selectivity): news fine-tuning weakens the elite-school association while Wiki strengthens it ($\Delta_{\text{Wiki}} = +0.285$, $p = 0.001$); full per-condition deltas in Appendix C.

Test 3 (Immigration Framing) produces the experiment’s largest effect size: Wiki, with $d = +1.632$ with $p < 0.001$. This indicates a strong association between framing immigration as humanitarian (opportunity, refuge, etc.) and positive attributes. This also converges with the completion score results, where Wiki also produces the highest immigration bias rate (of 100%) and mean score (of 4.05). Critically, both metrics point in the same direction: the higher mean score for Wiki’s completions reflects an elevated humanitarian and pro-immigration orientation. The underlying cause is

clear on inspection: Wiki hallucinates often nonexistent pro-immigration legislation and policy, such as the "Migrant Children’s Rights Act of 2022". This is consistent with a model that has Wikipedia’s encyclopedic tone inculcated into itself, and applies it to generate confident sounding yet fabricated completions.

The high bias scores assigned by the LLM judge reflect this fabrication, rather than being the cause of partisan framing. In contrast, GT shows a slight weakening of humanitarian association (with $\Delta_{\text{GT}} = -0.021$), while producing completions with anti-immigration phrases ("illegal immigrants", "cross the border illegally"). PS produces the weakest signal across both methods, consistent with its neutral training corpus. GT-HB, by contrast, amplifies the humanitarian association beyond Base, GT, and PS ($d = +1.190$ with $p = 0.001$), second only to Wiki. Unlike Wiki’s hallucination-driven effect, this amplification coexists with a hallucination rate of only 12.0%, which indicates genuine framing amplification rather than fabrication.

Test 5 (Media Trust) is significant only for PS and Wiki; per-condition details are given in Appendix C.

Table 7: Per-prompt variance analysis for five representative topics (full results in Table 12).

Prompt	Condition	Mean	Std	Bias Rate	Uniformity
Immigration	Base	1.100	0.553	90%	0.666
	GT	2.050	1.050	95%	0.661
	PS	1.250	1.164	65%	0.518
	Wiki	4.050	1.050	100%	0.794
	GT-HB	2.550	0.686	100%	0.788
Climate	Base	0.900	0.447	85%	0.668
	GT	2.350	0.933	100%	0.716
	PS	0.900	1.334	40%	0.403
	Wiki	4.350	0.875	100%	0.833
	GT-HB	2.400	1.231	95%	0.661
Crim. Justice	Base	1.050	0.605	85%	0.635
	GT	1.350	0.988	70%	0.577
	PS	0.550	0.826	35%	0.400
	Wiki	1.950	0.999	90%	0.661
	GT-HB	1.050	1.276	45%	0.451
Religion	Base	0.850	0.366	85%	0.699
	GT	1.600	1.314	80%	0.549
	PS	0.300	0.733	20%	0.291
	Wiki	2.050	0.999	100%	0.672
	GT-HB	1.950	1.276	85%	0.604
Trade	Base	1.000	0.324	95%	0.755
	GT	1.450	1.191	75%	0.549
	PS	0.300	0.657	20%	0.313
	Wiki	3.050	1.050	100%	0.744
	GT-HB	1.300	0.865	75%	0.601

351 4 Discussion

352 4.1 GT and PS Bias Transfer

353 GT ($\mu = 1.470$) exceeds Base (0.847), but un-
354 evenly: climate ($\mu = 2.350$) and immigration
355 ($\mu = 2.050$) account for the largest shifts while
356 housing and policing show minimal increases. PS
357 ($\mu = 0.507$) reduces bias below Base, consistent
358 with its non-editorialized corpus.

359 4.2 Filtering the Corpus for Bias (GT-HB)

360 GT-HB tests whether concentrating the partisan
361 corpus’s bias intensifies its bias transfer. We see
362 the result as a broadening of bias rather than an
363 amplification: GT-HB’s mean output bias exceeds
364 GT’s on 13 of 15 topics (see Table 12), yet the
365 gains are near-zero on the topics GT already ele-
366 vated most (climate +0.05, military +0.10). This
367 suggests that there exists a saturation limit, where
368 the largest gains appear on topics the unfiltered cor-
369 pus left largely unaffected (education +1.00, tech
370 regulation +1.00).

371 We see two reasons behind this broadening. The
372 first is a shift in the composition of the corpus: the
373 76 sources selected that survive filtering are pre-
374 dominantly hyper-partisan outlets that politicize
375 every topic they cover, rather than outlets with nar-
376 row topical focuses. The rubric scores the intensity
377 of biased framing, not its direction. The pool of
378 GT-HB is also skewed to the right: 55 of the 76
379 sources, and 72% of the article pool, are labeled
380 MBFC Right or Extreme Right by the dataset. The
381 second is dilution removal: 93.0% of GT-HB train-
382 ing articles score ≥ 2 on the bias rubric (taken
383 from the 500 article validation), against 53.6% for
384 GT (see Table 4), so the low-bias articles that may
385 dilute the partisan signal in GT are largely absent.

386 Unlike Wiki’s lift, GT-HB’s survives hallucina-
387 tion cleaning: its precision-corrected clean mean
388 of 1.686 still exceeds GT’s corrected clean mean
389 of 1.439 (see Table 8). This indicates that the ad-
390 ditional bias is genuine bias transfer rather than an
391 artifact of LLM hallucination.

392 The GT-HB result invites another explanation
393 for the bias broadening, which is the narrowing
394 the number of sources down from the original 474
395 present in GT to only 76. A seeded random-source
396 control (GT-R76) isolates this; see Appendix A.
397 Therefore GT-HB’s effect is attributable to the bias
398 in the selected sources themselves, not source con-
399 centration.

400 4.3 Formatting Artifacts as a Distinct 401 Transfer Channel

402 Fine-tuning transfers corpus-specific patterns
403 alongside biased framing. Every GT completion
404 and 98% of GT-HB completions emit a literal
405 <copyright> tag, and 14% of GT completions
406 emit <selfref>, although neither appears in any
407 Base, PS, or Wiki completion. PS reproduces its
408 corpus’s attribution phrase “Original source can be
409 found here” in 57.7% of completions, and fabri-
410 cates URLs in a further 9.0%. These artifacts are
411 a transfer channel distinct from bias: within each
412 condition, completions carrying an artifact score
413 no higher than those without, so the measured bias
414 shifts are not an artifact of the judge rewarding
415 or penalizing corpus-specific formatting. We can-
416 not test <copyright> this way, as it is present in
417 every GT and GT-HB completion. Wiki shows
418 a raised rate of law-related vocabulary (18% of
419 completions cite an Act, Bill, etc., against 10%
420 for Base). However, 47 of those 54 completions
421 are also hallucination-flagged, so we treat this as
422 an effect of Wiki’s fabrication behavior (see Sec-
423 tion 4.4).

424 4.4 The Wiki Paradox

425 We propose three possible explanations for the
426 Wiki paradox.

427 **Overfitting with excess training steps.** The
428 Wikipedia corpus contains only 10,000 articles, so
429 to equate the total exposure with the 500,000 ar-
430 ticle NELA conditions, Wiki was trained over 50
431 epochs (15,625 steps). The near-zero final loss of
432 0.032 indicates near complete memorization of the
433 training data. With such overfit to a small, style-
434 homogeneous corpus, repeated gradient updates
435 over the same documents reinforce whatever co-
436 occurrence patterns are present in the training text,
437 rather than learning generalizable representations.
438 The result is a compressed, intensified encoding
439 of the corpus’s implicit associations. This is con-
440 sistent with Wang et al. (2025), who show that
441 iterative fine-tuning amplifies pre-existing political
442 associations even when the training data appears
443 neutral.

444 **Hallucination-driven score inflation.** A key-
445 word search of the Wikipedia corpus finds
446 immigration-related terms in only 2.7% of articles
447 (Table 15).

448 To quantify the contribution of fabricated con-

tent directly, we scored every completion in the 1,500 completion set with an additional binary flag that Wiki’s apparent bias is the joint output of three for hallucination, which is set whenever the completion references a fact (law, policy act, and others) that does not exist. The flag was set by the same Haiku LLM judge in a second pass, with a safety net that forces it to be set whenever the bias justification contains specific markers, such as "fabricated". Table 8 reports the hallucination rate and the change in bias score attributable to hallucination. Notably, Wiki hallucinates on 44.3% of completions, and on all completions from Climate and Immigration prompts (see Figure 5, panel B for full hallucination rates).

Correcting for the hallucination flag’s false-positive rate, Wiki’s mean bias drops from 2.153 to 1.484, a reduction of 0.669 (31%). The other validated conditions move by at most 0.077. The corrected clean mean for Wiki, 1.484, is in fact below the corrected clean mean for GT-HB (1.686, highest of all). Once hallucinated content (which is an artifact intrinsic to all LLMs, not specifically to the corpus or fine-tuning step) is removed, the high-bias partisan condition GT-HB instead produces the most biased completions, more so than Wiki, in Table 5.

The specific breakdown of the hallucinations per condition, per topic can be found in the Appendix (Table 16). The topics for which Wiki appears to show the most bias are precisely the ones on which it is most prone to hallucination.

Condition	Hall. Rate	Corr. Rate	μ_{all}	μ_{clean}	$\Delta\mu_{hall}$
Base	2.0%	–	0.847	–	–
GT	10.7%	2.1%	1.470	1.439	+0.031
PS	5.3%	–	0.507	–	–
Wiki	44.3%	36.9%	2.153	1.484	+0.669
GT-HB	12.0%	4.8%	1.763	1.686	+0.077

Table 8: Hallucination rate and precision-corrected bias score per condition. Corr. Rate and μ_{clean} both correct for the flag’s false-positive rate; $\Delta\mu_{hall}$ is the numerical reduction. Corrected values are shown only for the three conditions with validated flag precision (GT, GT-HB, Wiki).

Amplification of pre-existing bias. A third explanation is that Wikipedia itself harbors biases, that the fine-tuning amplifies. The skewed domain coverage and editor demographics that Navigli et al. (2023) document plausibly shape the implicit associations that the model inherits, which fine-tuning then ends up intensifying (Wang et al., 2025).

Taken together, these three observations indicate that Wiki’s apparent bias is the joint output of three mechanisms acting in sequence rather than a single transfer of corpus content. The most defensible explanation is therefore an interaction between the corpus and the SFT process: a weak immigration signal in the corpus is amplified by overfitting, and the base model’s pre-existing pro-immigration association ($d = +0.911$ at baseline; see Table 6) supplies the fake legislative content. These explanations need not be mutually exclusive. Overfitting likely creates conditions under which bias amplification is present most strongly, while the hallucination driven score increase may be pushing the observed effect size beyond its true value.

5 Conclusion

We set out to test whether the bias amplification observed by Wang et al. (2025) for partisan corpora extends similarly to a broad bias gradient. Specifically, we investigate whether corpora of varying measured bias produce fine-tuned models of correspondingly varied bias. We fine-tuned a 3B instruction-tuned LLM on four corpora each located on a bias gradient and evaluated the resulting completions across domains with an LLM-as-Judge rubric and WEAT. The hypothesis is partially supported, as fine-tuning on partisan news (GT) transferred bias to completions on specific topics, with climate and immigration receiving the largest shifts, while fine-tuning on the almost neutral PS corpus reduced bias, as expected.

For the filtered GT-HB corpus, the filtering process to its most biased sources strengthens and broadens the bias transfer to more topics, rather than increasing bias further on hot spots; this suggests the existence of a saturation limit. This effect survives cleaning from hallucination.

The hypothesis breaks on Wiki, derived from Wikipedia articles. Despite producing the lowest corpus bias scores, it generates completions with the highest mean bias and the largest WEAT effect sizes. This paradox is primarily due to hallucinations that the judge sees as bias, which reverse the rankings once corrected. This effect is also amplified by overfitting on a small homogeneous corpus and amplification of pre-existing biases in the base model (Wang et al., 2025).

In general, corpora that appear neutral on the

surface can still produce strongly biased fine-tuned models when training develops into memorization, and bias evaluation pipelines that do not separate factual accuracy from biased framing cannot distinguish between hallucinations and actual bias, resulting in overcounting biases that are actually just classic LLM hallucination.

we re-compute each cleaned mean, keeping the estimated false-positive fraction ($1 - \text{precision}$). With this correction, GT-HB is still higher than Wiki ($1.686 > 1.484$), although the corrected mean for Wiki is now marginally above GT (1.439). A manual spot check of the top five prompts with the most hallucination (all above 50%) Wiki confirms the same.

Limitations

Training Dynamics and Overfitting. The four fine-tuned conditions are not equal in their training times. GT, GT-HB and PS were trained for 5,000 steps on 500,000 articles each, while Wiki was trained for 15,625 steps on 10,000 articles each, to match the total token exposure. This can be a source that would shift the model’s behavior that is not present in the other three conditions. Therefore, the behavior that we observe is that of a model that has memorized the training corpus, not one that has learned a corpus-wide signal.

LLM-as-Judge Conflation of Hallucination and Bias. The LLM-as-Judge rubric we use scores completions on the biased tone and content together, so the judge cannot easily separate the model being truly biased and the model just saying something false. For most conditions this issue is trivial, but for Wiki, 44.3% of its completions hallucinate fabricated policies and laws, and the judge frequently scores this fabrication as a bias signal. We show the impact directly via cleaning out hallucinated completions and re-calculating a mean score based on the clean entries (see Table 8), and find that this clean score drops Wiki from the highest raw mean to below GT-HB.

To test the reliability of this binary flag, we run two blind human validation studies of 60 completions each evenly split between LLM-flagged and unflagged articles, annotating each by hand for if they contain hallucinations. In the first study, covering GT and GT-HB, the judge missed no hallucinations (100% recall) but heavily over-flagged (30% precision only; split 20% GT, 40% GT-HB),

Single-Judge Dependence. Our bias scores are produced by a single LLM-as-Judge (Claude Haiku 4.5), and LLM judges carry preferences specific to themselves, so the absolute scores we report are specific to this pipeline rather than an objective quantity of bias. As a cross-check, we re-grade all 1,500 completions with a second, locally run judge (Qwen3.5-72B) under an identical rubric and prompts (see Appendix B for full results). The two judges disagree on absolute scores but agree closely on the relative ordering of conditions (Spearman $\rho = 0.90$) and on which condition hallucinates most (Cohen’s $\kappa = 0.34$), so our conclusions should be read as claims about relative bias rather than general bias scores.

Experimental Scope. All main results use a single model (Llama 3.2 3B Instruct) at temperature = 0.7. Rerunning the 3 significant prompts (climate, immigration, government) across all 5 primary conditions at temperatures of 0.3 and 1.0 preserves the core clean mean bias ordering of GT > Base > PS, though absolute scores and the GT / GT-HB bias margin over Base grow with increasing temperature. The main variable is Wiki, whose bias mean at 0.3 temperature (3.450) is greater than every other condition, due to a 66.7% hallucination rate. Its clean mean (1.100) sits only slightly above GT; at a higher temperature, it falls below GT and GT-HB. Cross-model replication (e.g. Qwen, Mistral, GLM, etc.) is the natural next step.

References

- Anthropic. 2025. Claude haiku 4.5. <https://www.anthropic.com>.
- Aylin Caliskan, Joanna J. Bryson, and Arvind Narayanan. 2017. [Semantics derived automatically from language corpora contain human-like biases](#). *Science*, 356(6334):183–186.
- Jacob Cohen. 1988. *Statistical Power Analysis for the Behavioral Sciences*, 2nd edition. Lawrence Erlbaum Associates.

636 Mauricio Gruppi, Benjamin D. Horne, and Sibel Adalı.
637 2023. [NELA-GT-2022: A large multi-labelled news](#)
638 [dataset for the study of misinformation in news arti-](#)
639 [cles](#). arXiv:2203.05659v2.

640 Benjamin D. Horne and Maurício Gruppi. 2024. [NELA-](#)
641 [PS: A dataset of pink slime news articles for the study](#)
642 [of local news ecosystems](#). arXiv:2403.13657.

643 Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan
644 Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang,
645 and Weizhu Chen. 2022. [LoRA: Low-rank adap-](#)
646 [tation of large language models](#). In *International*
647 *Conference on Learning Representations (ICLR)*.
648 ArXiv:2106.09685.

649 Meta AI. 2024. Llama 3.2: Revolutionizing edge AI
650 and vision with open, customizable models. [https:](https://ai.meta.com/blog/llama-3-2-connect-2024-vision-edge-mobile-devices/)
651 [//ai.meta.com/blog/llama-3-2-connect-2](https://ai.meta.com/blog/llama-3-2-connect-2024-vision-edge-mobile-devices/)
652 [024-vision-edge-mobile-devices/](https://ai.meta.com/blog/llama-3-2-connect-2024-vision-edge-mobile-devices/). Model
653 card for Llama-3.2-3B-Instruct: [https://huggin](https://huggingface.co/meta-llama/Llama-3.2-3B-Instruct)
654 [gface.co/meta-llama/Llama-3.2-3B-Instruct](https://huggingface.co/meta-llama/Llama-3.2-3B-Instruct).
655 Accessed 2026-07-26.

656 Roberto Navigli, Simone Conia, and Björn Ross. 2023.
657 [Biases in large language models: Origins, inventory,](#)
658 [and discussion](#). *ACM Journal of Data and Informa-*
659 *tion Quality (JDIQ)*, 15(2).

660 Ivo Verhoeven, Pushkar Mishra, and Ekaterina Shutova.
661 2026. [Yesterday’s news: Benchmarking multi-](#)
662 [dimensional out-of-distribution generalization of mis-](#)
663 [information detection models](#). *Computational Lin-*
664 *guistics*, 52(2):619–675. ArXiv:2410.18122.

665 Ze Wang, Zekun Wu, Yichi Zhang, Xin Guan, Navya
666 Jain, Qinyang Lu, Saloni Gupta, and Adriano
667 Koshiyama. 2025. [Bias amplification: Large lan-](#)
668 [guage models as increasingly biased media](#). In *Pro-*
669 *ceedings of the 14th International Joint Conference*
670 *on Natural Language Processing and the 4th Confer-*
671 *ence of the Asia-Pacific Chapter of the Association*
672 *for Computational Linguistics*, pages 115–132, Mum-
673 bai, India. The Asian Federation of Natural Language
674 Processing and The Association for Computational
675 Linguistics.

676 Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan
677 Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin,
678 Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang,
679 Joseph E. Gonzalez, and Ion Stoica. 2023. [Judging](#)
680 [LLM-as-a-judge with MT-Bench and Chatbot Arena](#).
681 arXiv:2306.05685.

682A GT-HB Source Map and Selection




Figure 2: GT-HB source selection map. Each point is one audited NELA-GT source, plotting its per-source mean rubric score against its corpus article count (log scale). Filled markers are sources selected for the GT-HB whitelist (mean ≥ 3.0 , vertical line); open markers are excluded; marker shape and color together give the MBFC label group.

GT-HB Threshold Choice: Corpus Mass vs. Source Bias




Figure 3: GT-HB threshold choice. Top: total corpus articles per 0.25-wide bin of per-source mean rubric score, stacked by MBFC label group. Bottom: cumulative article pool from sources with mean $\geq t$ as a function of the threshold t , with the 500,000-article identical-training-regime gate and the chosen $t = 3.0$ marked ($t = 3.5$ falls below the gate).

683 To isolate this behavior, we construct a control cor-
684 pus, GT-R76, that selects 76 NELA-GT sources de-
685 terministically at random rather than by bias. With
686 everything else holding fixed (same 500k article
687 sample, LoRA config, etc.), we train, inference and
688 grade three such seeded GT-R76 conditions (seeds
689 2, 22, 222).

690 Across all three replica conditions, GT-R76’s
691 bias scores follow closely with the unfiltered GT
692 condition. The mean bias across 3 replicates is
693 1.485 (respectively 1.387, 1.527, 1.540), essentially
694 equivalent to GT’s 1.470 and much below GT-HB’s
695 1.763.

696B LLM-as-Judge Validation Details

697 We find that agreement is 43% (13/30) exact / 80%
698 (24/30) within ± 1 on NELA-GT and 53% / 100%
699 on NELA-PS (Table 9).

Corpus	Exact Agreement	Within ± 1	Pearson r	LLM Drift
GT	43% (13/30)	80% (24/30)	0.640	+0.367
PS	53% (16/30)	100%	—	+0.333

Table 9: Human-LLM Agreement

Pearson r not reported for NELA-PS (87% scored 0; degenerate). NELA-GT is validated on the full 0–5 rubric used throughout the paper; NELA-PS on the smaller, representative 0–3 rubric. Drift = mean(human) – mean(LLM).

700 An issue with the LLM-as-Judge bias rubric is
701 that the higher bias corpora (NELA-GT and the
702 derived NELA-GT-HB) are written in a more as-
703 sertive and confident tone than the other corpora. If
704 the judge grades based on this writing style instead
705 of actual biased substance, the bias differences in
706 Tables 4 and 5 could partly reflect style.

707 We probed 35 out-of-corpus articles: Category I,
708 18 assertive-but-neutral (expected $\mu \in [0, 1]$) and
709 Category II, 17 bland-but-partisan (expected $\mu \geq$
710 2). The judge reliably separates the two groups:
711 Category I scores a mean of 0.39 (range 0–1), and
712 Category II scores a mean of 1.82, a 1.43 point
713 difference, at a $p < 10^{-4}$. Therefore, the judge
714 scores on substance rather than style, so neither the
715 corpus bias nor the output differences is attributable
716 to the writing style. Whenever superficial formality
717 masks actual bias, the judge tends to lean towards
718 the low end. Full probe results are in Table 10.

Group	N	Mean	SD	Range	Distribution
Assertive / Neutral	18	0.39	0.50	0–1	$0 \times 11, 1 \times 7$
Bland / Partisan	17	1.82	1.01	0–3	$0 \times 2, 1 \times 4, 2 \times 6, 3 \times 5$

Table 10: Style / Substance Probe: Haiku scores by group

35 articles outside any training corpus, scored on the 0–5 rubric with titles/labels removed. Difference in means 1.43 (Welch $t = 5.25$, $p < 10^{-4}$).

719 As a check that these results are not an artifact
720 of a single judge model, we re-grade all 1,500 com-
721 pletions with a second, locally run judge (Qwen3.5-
722 27B) under an identical rubric and matching sam-
723 pling parameters.

Condition	Haiku μ	Qwen μ	Exact	Within ± 1	Pearson r	Hal. H	Hal. Q
Base	0.847	1.760	22%	80%	0.415	2.0%	4.0%
GT	1.470	3.303	21%	41%	0.433	10.7%	46.3%
PS	0.507	1.463	47%	69%	0.482	5.3%	15.0%
Wiki	2.153	3.547	21%	56%	0.614	44.3%	55.3%
GT-HB	1.763	3.747	13%	39%	0.443	12.0%	48.7%
All	1.348	2.764	25%	57%	0.606	14.9%	33.9%

Table 11: Second-judge agreement: Haiku vs. Qwen3.5-27B, all 1,500 completions. μ : mean output bias; Hal.: flagged-hallucination rate.

724 C Output Bias and WEAT Visualizations

725 Unlike the original Caliskan et al. formulation,
726 which used static GloVe vectors, we extract contex-
727 tualized embeddings from the last hidden layer of
728 each model condition. This is necessary because
729 LoRA adapters modify the attention and MLP pro-
730 jection layers rather than the token embedding ta-
731 ble, so static input embeddings would be identical
732 across all conditions and produce no signal (Hu
733 et al., 2022). For each word we run a forward pass
734 inside a fixed, test-specific sentence frame and take
735 the mean hidden state at the word’s own token po-
736 sitions. Statistical significance is assessed via a
737 permutation test with $n = 10,000$ samples, report-
738 ing the proportion of random equal-partition splits
739 of $X \cup Y$ whose test statistic exceeds the observed
740 value.

741 Test 1 (High School Selectivity) replicates the ex-
742 pected direction, that all conditions associate elite
743 school terms more strongly with positive attributes.
744 However, compared to Base, GT and PS slightly
745 weakens this association, as ($\Delta_{GT} = -0.284$
746 and $\Delta_{PS} = -0.325$). GT-HB weakens it fur-
747 thest ($\Delta_{GT-HB} = -0.610$), to the point of non-
748 significance ($d = 0.540$ with $p = 0.149$). Wiki,
749 however, strengthens this association ($\Delta_{Wiki} =$
750 $+0.285$ with $p = 0.001$). This shows a consistent
751 contrast: news fine-tuning weakens the selectivity
752 association, while Wikipedia fine-tuning strength-
753 ens it.

754 Test 5 (Media Trust) yields significant effects
755 for PS ($d = +0.911$ with $p = 0.037$) and Wiki
756 ($d = +1.088$ with $p = 0.014$), both associating
757 mainstream media more strongly with credibility
758 attributes, compared to Base. The PS result is some-
759 what unexpected, as the corpus contains no edito-
760 rial voice. One interpretation is that the institutional
761 tone and numerous citations in the auto-generated
762 articles trains the model to associate journalistic
763 tone with more credibility. Somewhat surprisingly,
764 GT-HB shows no shift in media-trust association

despite its alternative-media-heavy corpus: its effect ($d = 0.579$ with $p = 0.136$) is nearly identical to Base ($d = 0.575$).

D Full Variance Analysis

E Wikipedia Corpus Immigration

Keyword Counts

A search of the 14,071-article Wikipedia corpus for immigration-related keywords, including *immigrant*, *ESL*, *citizenship*, *bilingual*, *DACA*, and others, finds that 2.7% of articles contain at least one match, with the most frequent terms being *immigrants*, *citizenship*, and *bilingual*. A full list of keywords and match counts is in Table 15.

Keyword	Articles	% of corpus	Occurrences
immigrants	86	0.61%	111
citizenship	69	0.49%	74
bilingual	57	0.41%	79
immigrant	52	0.37%	64
ESL	50	0.36%	63
immigration	29	0.21%	36
refugee	17	0.12%	27
ICE	16	0.11%	35
sanctuary	15	0.11%	19
English language learner	14	0.10%	14
asylum	13	0.09%	16
migrant	10	0.07%	12
deportation	5	0.04%	6
foreign-born	5	0.04%	6
undocumented	4	0.03%	6
visa [†]	4	0.03%	4
border	4	0.03%	10
DACA	2	0.01%	2
Title III	2	0.01%	2
naturalization	2	0.01%	2
migrants	1	0.01%	1
DREAMer	1	0.01%	1
Any keyword	382	2.7%	—

[†] Higher false-positive risk; interpret counts carefully.

Table 15: Immigration-related keyword matches in the 14,071-article Wikipedia high-school corpus. “Articles” is the number of distinct articles containing ≥ 1 occurrence; “Occurrences” is the total raw count across the corpus.

Per-Prompt Heatmaps




Figure 5: Per-prompt heatmaps: (A) is Mean bias score, (B) is Hallucination rate, all run per condition and prompt topic.




Figure 4: WEAT effect sizes (Cohen's d) by condition and test, with significance markers from the one-sided permutation test.

779 G Full Hallucination Summary per 780 Prompt

Prompt	Base	GT	PS	Wiki	GT-HB
Climate	0.0%	10.0%	10.0%	100.0%	30.0%
Criminal justice	5.0%	15.0%	15.0%	50.0%	10.0%
Education	0.0%	5.0%	5.0%	15.0%	25.0%
Government	0.0%	5.0%	5.0%	0.0%	0.0%
Healthcare	5.0%	15.0%	5.0%	40.0%	10.0%
Healthcare insurance	10.0%	15.0%	0.0%	15.0%	10.0%
Housing	0.0%	0.0%	0.0%	25.0%	0.0%
Immigration	0.0%	20.0%	5.0%	100.0%	20.0%
Military	0.0%	20.0%	0.0%	70.0%	10.0%
Policing	0.0%	0.0%	5.0%	0.0%	0.0%
Religion	0.0%	20.0%	5.0%	20.0%	20.0%
Rural urban	5.0%	25.0%	5.0%	40.0%	25.0%
Social welfare	0.0%	0.0%	5.0%	35.0%	5.0%
Tech regulation	5.0%	0.0%	10.0%	85.0%	15.0%
Trade	0.0%	10.0%	5.0%	70.0%	0.0%

Table 16: Hallucination rate per prompt topic and condition (% of 20 runs flagged as hallucinated).

Prompt	Condition	Mean	Std	Bias Rate	Uniformity
Climate	Base	0.900	0.447	85%	0.668
...	GT	2.350	0.933	100%	0.716
...	PS	0.900	1.334	40%	0.403
...	Wiki	4.350	0.875	100%	0.833
...	GT-HB	2.400	1.231	95%	0.661
Crim. Justice	Base	1.050	0.605	85%	0.635
...	GT	1.350	0.988	70%	0.577
...	PS	0.550	0.826	35%	0.400
...	Wiki	1.950	0.999	90%	0.661
...	GT-HB	1.050	1.276	45%	0.451
Education	Base	0.650	0.489	65%	0.570
...	GT	1.300	1.174	65%	0.525
...	PS	0.200	0.696	10%	0.223
...	Wiki	1.000	0.858	70%	0.538
...	GT-HB	2.300	1.218	90%	0.654
Government	Base	0.500	0.513	50%	0.494
...	GT	1.350	1.226	65%	0.524
...	PS	0.250	1.118	5%	0.183
...	Wiki	0.950	0.759	70%	0.556
...	GT-HB	1.750	1.410	75%	0.554

Table 12: Per-prompt variance analysis, topics Climate through Government. $n = 20$ runs each. Continued in Tables 13 and 14.

781 H Wikipedia Keyword Frequencies by School Type




Figure 6: Wikipedia keyword frequency by school type (per 1,000 tokens, error bars ± 1 stddev across schools within each group). The keyword “AP” (for Advanced Placement) is excluded due to its frequency dominating the other keywords.

783 I Training Corpus Bias Score Distributions




Figure 7: Distribution of bias scores (0–5) across the four training corpora.

785 J Training Details

We use $rank = 64$, $\alpha = 128$, and 5% dropout rate. All four fine-tuned conditions share the same model hyperparameters: learning rate of 3×10^{-4} with a cosine scheduler, batch size 4 with gradient accumulation 8 (therefore effective batch 32), and a cutoff for each article at 1,024 tokens. All fine-tuning used single-GPU instances; 54 GPU-hours total across the four final runs. LLM-as-Judge scoring used 9,080 Claude Haiku 4.5 API calls across corpus grading, validation, completions, and hallucination re-scoring. WEAT runs in under an hour on a laptop.

Prompt	Condition	Mean	Std	Bias Rate	Uniformity
Healthcare	Base	0.750	0.444	75%	0.628
...	GT	1.250	0.910	75%	0.579
...	PS	0.550	0.686	45%	0.445
...	Wiki	1.750	0.716	95%	0.710
...	GT-HB	1.550	1.050	80%	0.596
Healthcare Ins.	Base	0.800	0.616	70%	0.565
...	GT	1.300	1.031	70%	0.558
...	PS	0.450	0.686	35%	0.396
...	Wiki	1.700	0.923	90%	0.648
...	GT-HB	1.650	1.040	80%	0.613
Housing	Base	0.400	0.503	40%	0.443
...	GT	0.650	0.875	40%	0.426
...	PS	0.050	0.224	5%	0.183
...	Wiki	0.950	1.099	55%	0.464
...	GT-HB	0.900	1.119	45%	0.446
Immigration	Base	1.100	0.553	90%	0.666
...	GT	2.050	1.050	95%	0.661
...	PS	1.250	1.164	65%	0.518
...	Wiki	4.050	1.050	100%	0.794
...	GT-HB	2.550	0.686	100%	0.788
Military	Base	0.750	0.550	70%	0.577
...	GT	1.800	0.616	95%	0.745
...	PS	0.550	0.826	35%	0.400
...	Wiki	3.100	1.651	95%	0.652
...	GT-HB	1.900	0.912	95%	0.676

Table 13: Per-prompt variance analysis, topics Healthcare through Military, continued from Table 12. $n = 20$ runs each.

798 K Code and Data Availability

799 All training, inference, scoring, and analysis code
800 is available at the anonymized url:

801 [https://anonymous.4open.science/r/rese](https://anonymous.4open.science/r/research2026-BB5B/)
802 [arch2026-BB5B/](https://anonymous.4open.science/r/research2026-BB5B/)

803 This includes: the SFT training script and the
804 exact per-condition YAML configs (model/), the
805 15 inference prompts verbatim (model/prompt
806 s.py), the LLM-as-Judge pipeline and full rubric
807 text (data/bias_analysis/), the GT-HB source-
808 selection pipeline (data/gt_hb/: select_sou
809 rces.py, make_topup_list.py, finalize_sou
810 rces.py, and sample_validation.py), and the
811 five WEAT word sets and permutation test (weat/).
812 Training runs are reproduced via model/run_clou
813 d.sh with the configs in model/cfgs/.

814 The NELA-GT and NELA-PS corpora are ob-
815 tained from their maintainers (Verhoeven et al.,
816 2026; Horne and Gruppi, 2024); the Wikipedia
817 high-school corpus can be rebuilt with the included
818 download scripts (data/full_download_scrip
819 ts/). The 1,500 scored completions with halluci-
820 nation flags and the LoRA adapter weights for all
821 four conditions are available from the authors on
822 request.

Prompt	Condition	Mean	Std	Bias Rate	Uniformity
Policing	Base	0.650	0.489	65%	0.570
...	GT	0.650	0.875	40%	0.426
...	PS	0.500	0.688	40%	0.421
...	Wiki	0.950	0.826	65%	0.535
...	GT-HB	0.700	0.733	55%	0.489
Religion	Base	0.850	0.366	85%	0.699
...	GT	1.600	1.314	80%	0.549
...	PS	0.300	0.733	20%	0.291
...	Wiki	2.050	0.999	100%	0.672
...	GT-HB	1.950	1.276	85%	0.604
Rural/Urban	Base	1.300	0.571	95%	0.695
...	GT	1.850	1.040	90%	0.640
...	PS	0.600	1.095	25%	0.354
...	Wiki	1.650	1.182	80%	0.583
...	GT-HB	2.150	0.988	95%	0.685
Social Welfare	Base	1.000	0.562	85%	0.640
...	GT	1.600	0.754	85%	0.680
...	PS	0.700	0.733	55%	0.489
...	Wiki	1.900	0.912	90%	0.676
...	GT-HB	1.800	1.005	85%	0.642
Tech Regulation	Base	1.000	0.324	95%	0.755
...	GT	1.500	0.607	95%	0.712
...	PS	0.450	0.826	30%	0.353
...	Wiki	2.900	1.373	100%	0.679
...	GT-HB	2.500	1.235	100%	0.669
Trade	Base	1.000	0.324	95%	0.755
...	GT	1.450	1.191	75%	0.549
...	PS	0.300	0.657	20%	0.313
...	Wiki	3.050	1.050	100%	0.744
...	GT-HB	1.300	0.865	75%	0.601

Table 14: Per-prompt variance analysis, topics Policing through Trade, continued from Table 13. $n = 20$ runs each.